



Housing Price Prediction

Submitted by:
Abhishek Kumar

ACKNOWLEDGMENT

I would like to thanks to Flip Robo Technologies to give me a wonderful opportunity. This project is given by my SME Ms Sapna Verma. I have referred below resources that helped and guided me in completion of this project as below:-

- <https://www.kaggle.com/erick5/predicting-house-prices-with-machine-learning>
- <https://studygyaan.com/data-science-ml/linear-regression-machine-learning-project-for-house-price-prediction>
- https://loddonhouse.co.uk/?gclid=CjwKCAjw-ZCKBhBkEiwAM4qfF_ZWhedS9VWDcP3TZ5_SVB7xuurHYsU5s4MaQzoRhiVB5fnbA1I1DxoC3G0QAvD_BwE

INTRODUCTION

Business Problem Framing

- Houses are one of the necessary need of each and every person around the globe and therefore housing and real estate market is one of the markets which is one of the major contributors in the world's economy. It is a very large market and there are various companies working in the domain. Data science comes as a very important tool to solve problems in the domain to help the companies increase their overall revenue, profits, improving their marketing strategies and focusing on changing trends in house sales and purchases. Predictive modelling, Market mix modelling, recommendation systems are some of the machine learning techniques used for achieving the business goals for housing companies. Our problem is related to one such housing company.
- A US-based housing company named Surprise Housing has decided to enter the Australian market. The company uses data analytics to purchase houses at a price below their actual values and flip them at a higher price. For the same purpose, the company has collected a data set from the sale of houses in Australia. The data is provided in the CSV file below.
- The company is looking at prospective properties to buy houses to enter the market. You are required to build a model using Machine Learning in order to predict the actual value of the prospective properties and decide whether to invest in them or not. For this company wants to know:
 - Which variables are important to predict the price of variable?
 - How do these variables describe the price of the house?

Conceptual Background of the Domain Problem

Predicting sale prices for houses, even stranger ones. Use a test-driven approach to build a Linear Regression model using Python from scratch. We will use our trained model to predict house sale prices and extend it to a multivariate Linear Regression.

Review of Literature

We are required to model the price of houses with the available independent variables.

Technical Requirements:

- Data contains 1460 entries each having 81 variables.
- Data contains Null values. We need to treat them using the domain knowledge and your own understanding.
- Extensive EDA has to be performed to gain relationships of important variable and price.
- Data contains numerical as well as categorical variable. We need to handle them accordingly.
- We have to build Machine Learning models, apply regularization and determine the optimal values of Hyper Parameters.
- We need to find important features which affect the price positively or negatively.

- Two datasets are being provided to us (test.csv, train.csv). We will train on train.csv dataset and predict on test.csv file.

The "Data file.csv" and "Data description.txt" are enclosed with this file.

Motivation for the Problem Undertaken

- House is one of the important elements in basic human needs. People need a house to stay away from danger, hot weather, rainy day and as well as a place to stay calm. As long as people can fill the comfort of living under a roof then it is called a house. However, the things that matter is that the affordability of a person to purchase a house. Some people can afford a house that is really comfortable to stay in and some not. People who are called the rich and famous can afford a house that is almost called a heaven and some can only live in an ordinary but comfortable house. But it doesn't matter how our house may look like because the price of house is what that matter. We can see that the housing price is increasing as the time goes by. This may be an important area to look upon because more or less it could affect the economic level of a country. Therefore, a housing price can be defined as the rate of payment that one has to pay in order to purchase a house and for sure there are several factors that lead to housing price determination.
- In my own point of view, I believe that the increment of a housing price is due to the price increment in the raw material. Many may have similar idea but after looking into 10 journals as references for my proposed topic, I have found out several more important variables that lead to the factors of housing price determination. There are few number of knowledgeable individual turned up and able to find the contributing factors in determination of housing price. One who has studied using an empirical analysis has shown that income (demography trends) and nominal interest rates are the key explanatory factors in housing price. On the other hand, the equity returns may also have been an influential factor in the determination of housing price.

Analytical Problem Framing

Mathematical/ Analytical Modeling of the Problem

- This problem is a Linear Regression problem. The dataset is in CSV format and It contains 1460 training data points and 81 features that might help us predict the selling price of a house.
- Build a model of housing prices to predict median house values in California using the provided dataset.
- Train the model to learn from the data to predict the median housing price in any district, given all the other metrics.
- Predict housing prices based on median_income and plot the regression chart for it.

Data Sources and their formats

This Dataset is provided by Flip Robo Technologies CSV format. In this dataset, there are 1460 rows and 81 columns.

Load Data

```
In [92]: #uploading test dataset
test=pd.read_csv("Housing_test.csv")
train=pd.read_csv("Housing_train.csv")
```

In [93]: test

```
Out[93]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	LotConfig	LandSlope	Neighborhood	Condition1	Condition2
0	337	20	RL	88.0	14157	Pave	NaN	IR1	HLS	AllPub	Corner	Gtl	StoneBr	Norm	Norm
1	1018	120	RL	NaN	5814	Pave	NaN	IR1	Lvl	AllPub	CulDSac	Gtl	StoneBr	Norm	Norm
2	929	20	RL	NaN	11838	Pave	NaN	Reg	Lvl	AllPub	Inside	Gtl	CollgCr	Norm	Norm
3	1148	70	RL	75.0	12000	Pave	NaN	Reg	Bnk	AllPub	Inside	Gtl	Crawfor	Norm	Norm
4	1227	80	RL	88.0	14598	Pave	NaN	IR1	Lvl	AllPub	CulDSac	Gtl	Somerst	Feedr	Feedr
...
287	83	20	RL	78.0	10206	Pave	NaN	Reg	Lvl	AllPub	Inside	Gtl	Somerst	Norm	Norm
288	1048	20	RL	57.0	9245	Pave	NaN	IR2	Lvl	AllPub	Inside	Gtl	CollgCr	Norm	Norm
289	17	20	RL	NaN	11241	Pave	NaN	IR1	Lvl	AllPub	CulDSac	Gtl	NAmes	Norm	Norm
290	523	50	RM	50.0	5000	Pave	NaN	Reg	Lvl	AllPub	Corner	Gtl	BrkSide	Feedr	Feedr
291	1379	180	RM	21.0	1953	Pave	NaN	Reg	Lvl	AllPub	Inside	Gtl	BrDale	Norm	Norm

292 rows x 16 columns

In [94]: train

```
Out[94]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	LotConfig	LandSlope	Neighborhood	Condition1	Condition2
0	127	120	RL	NaN	4928	Pave	NaN	IR1	Lvl	AllPub	Inside	Gtl	NPkVill	Norm	Norm
1	889	20	RL	95.0	15985	Pave	NaN	IR1	Lvl	AllPub	Inside	Mod	NAmes	Norm	Norm
2	793	80	RL	92.0	9920	Pave	NaN	IR1	Lvl	AllPub	CulDSac	Gtl	NoRidge	Norm	Norm
3	110	20	RL	105.0	11751	Pave	NaN	IR1	Lvl	AllPub	Inside	Gtl	NWAmes	Norm	Norm
4	422	20	RL	NaN	16635	Pave	NaN	IR1	Lvl	AllPub	FR2	Gtl	NWAmes	Norm	Norm
...
1163	289	20	RL	NaN	9819	Pave	NaN	IR1	Lvl	AllPub	Inside	Gtl	Sawyer	Norm	Norm
1164	554	20	RL	87.0	8777	Pave	NaN	Reg	Lvl	AllPub	Inside	Gtl	Edwards	Feedr	Feedr
1165	198	180	RL	24.0	2280	Pave	NaN	Reg	Lvl	AllPub	FR2	Gtl	NPkVill	Norm	Norm
1166	34	70	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPub	Inside	Gtl	NPkVill	Feedr	Feedr

```
In [102]: #letscheck columns name of both dataset
print(train.columns)
print("*****")
print(test.columns)
```

```
Index(['MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea', 'Street', 'Alley',
       'LotShape', 'LandContour', 'Utilities', 'LotConfig', 'LandSlope',
       'Neighborhood', 'Condition1', 'Condition2', 'BldgType', 'HouseStyle',
       'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd', 'RoofStyle',
       'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType', 'MasVnrArea',
       'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual', 'BsmtCond',
       'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1', 'BsmtFinType2',
       'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', 'Heating', 'HeatingQC',
       'CentralAir', 'Electrical', '1stFlrSF', '2ndFlrSF', 'LowQualFinSF',
       'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath',
       'BedroomAbvGr', 'KitchenAbvGr', 'KitchenQual', 'TotRmsAbvGrd',
       'Functional', 'Fireplaces', 'FireplaceQu', 'GarageType', 'GarageYrBlt',
       'GarageFinish', 'GarageCars', 'GarageArea', 'GarageQual', 'GarageCond',
       'PavedDrive', 'WoodDeckSF', 'OpenPorchSF', 'EnclosedPorch', '3SsnPorch',
       'ScreenPorch', 'PoolArea', 'PoolQC', 'Fence', 'MiscFeature', 'MiscVal',
       'MoSold', 'YrSold', 'SaleType', 'SaleCondition', 'SalePrice'],
      dtype='object')
*****
Index(['MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea', 'Street', 'Alley',
       'LotShape', 'LandContour', 'Utilities', 'LotConfig', 'LandSlope',
       'Neighborhood', 'Condition1', 'Condition2', 'BldgType', 'HouseStyle',
       'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd', 'RoofStyle',
       'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType', 'MasVnrArea',
       'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual', 'BsmtCond',
       'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1', 'BsmtFinType2',
       'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', 'Heating', 'HeatingQC',
       'CentralAir', 'Electrical', '1stFlrSF', '2ndFlrSF', 'LowQualFinSF',
       'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath',
       'BedroomAbvGr', 'KitchenAbvGr', 'KitchenQual', 'TotRmsAbvGrd',
       'Functional', 'Fireplaces', 'FireplaceQu', 'GarageType', 'GarageYrBlt',
       'GarageFinish', 'GarageCars', 'GarageArea', 'GarageQual', 'GarageCond',
       'PavedDrive', 'WoodDeckSF', 'OpenPorchSF', 'EnclosedPorch', '3SsnPorch',
       'ScreenPorch', 'PoolArea', 'PoolQC', 'Fence', 'MiscFeature', 'MiscVal',
       'MoSold', 'YrSold', 'SaleType', 'SaleCondition'],
      dtype='object')
```

```
#check information of train and test dataset to find null values and type of columns
print(test.info())
print('*****')
print(train.info())
```

```
In [103]: #check information of train and test dataset to find null values and type of columns
print(test.info())
print('*****')
print(train.info())

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 292 entries, 0 to 291
Data columns (total 79 columns):
#   Column          Non-Null Count  Dtype
---  -
0   MSSubClass       292 non-null   int64
1   MSZoning         292 non-null   object
2   LotFrontage     247 non-null   float64
3   LotArea         292 non-null   int64
4   Street          292 non-null   object
5   Alley           14 non-null    object
6   LotShape        292 non-null   object
7   LandContour     292 non-null   object
8   Utilities       292 non-null   object
9   LotConfig       292 non-null   object
10  LandSlope       292 non-null   object
11  Neighborhood     292 non-null   object
12  Condition1      292 non-null   object
13  Condition2      292 non-null   object
14  BldgType        292 non-null   object
15  HouseStyle      292 non-null   object
16  OverallQual     292 non-null   int64
17  OverallCond     292 non-null   int64
18  YearBuilt       292 non-null   int64
19  YearRemodAdd    292 non-null   int64
20  RoofStyle       292 non-null   object
21  RoofMatl        292 non-null   object
22  Exterior1st     292 non-null   object
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 292 entries, 0 to 291
Data columns (total 79 columns):
#   Column          Non-Null Count  Dtype
---  -
0   MSSubClass       292 non-null   int64
1   MSZoning         292 non-null   object
2   LotFrontage     247 non-null   float64
3   LotArea         292 non-null   int64
4   Street          292 non-null   object
5   Alley           14 non-null    object
6   LotShape        292 non-null   object
7   LandContour     292 non-null   object
8   Utilities       292 non-null   object
9   LotConfig       292 non-null   object
10  LandSlope       292 non-null   object
11  Neighborhood     292 non-null   object
12  Condition1      292 non-null   object
13  Condition2      292 non-null   object
14  BldgType        292 non-null   object
15  HouseStyle      292 non-null   object
16  OverallQual     292 non-null   int64
17  OverallCond     292 non-null   int64
18  YearBuilt       292 non-null   int64
19  YearRemodAdd    292 non-null   int64
20  RoofStyle       292 non-null   object
21  RoofMatl        292 non-null   object
22  Exterior1st     292 non-null   object
```

23	Exterior2nd	292 non-null	object
24	MasVnrType	291 non-null	object
25	MasVnrArea	291 non-null	float64
26	ExterQual	292 non-null	object
27	ExterCond	292 non-null	object
28	Foundation	292 non-null	object
29	BsmtQual	285 non-null	object
30	BsmtCond	285 non-null	object
31	BsmtExposure	285 non-null	object
32	BsmtFinType1	285 non-null	object
33	BsmtFinSF1	292 non-null	int64
34	BsmtFinType2	285 non-null	object
35	BsmtFinSF2	292 non-null	int64
36	BsmtUnfSF	292 non-null	int64
37	TotalBsmtSF	292 non-null	int64
38	Heating	292 non-null	object
39	HeatingQC	292 non-null	object
40	CentralAir	292 non-null	object
41	Electrical	291 non-null	object
42	1stFlrSF	292 non-null	int64
43	2ndFlrSF	292 non-null	int64
44	LowQualFinSF	292 non-null	int64
45	GrLivArea	292 non-null	int64
46	BsmtFullBath	292 non-null	int64
47	BsmtHalfBath	292 non-null	int64
48	FullBath	292 non-null	int64
49	HalfBath	292 non-null	int64
50	BedroomAbvGr	292 non-null	int64
51	KitchenAbvGr	292 non-null	int64
52	KitchenQual	292 non-null	object
53	TotRmsAbvGrd	292 non-null	int64
54	Functional	292 non-null	object
55	Fireplaces	292 non-null	int64
56	FireplaceQu	153 non-null	object
57	GarageType	275 non-null	object
58	GarageYrBlt	275 non-null	float64
59	GarageFinish	275 non-null	object
60	GarageCars	292 non-null	int64
61	GarageArea	292 non-null	int64
62	GarageQual	275 non-null	object
63	GarageCond	275 non-null	object
64	PavedDrive	292 non-null	object
65	WoodDeckSF	292 non-null	int64
66	OpenPorchSF	292 non-null	int64
67	EnclosedPorch	292 non-null	int64
68	3SsnPorch	292 non-null	int64
69	ScreenPorch	292 non-null	int64

```

70 PoolArea      292 non-null  int64
71 PoolQC        0 non-null   float64
72 Fence         44 non-null   object
73 MiscFeature    10 non-null   object
74 MiscVal        292 non-null  int64
75 MoSold         292 non-null  int64
76 YrSold         292 non-null  int64
77 SaleType       292 non-null  object
78 SaleCondition  292 non-null  object
dtypes: float64(4), int64(33), object(42)
memory usage: 180.3+ KB

```

None

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 1168 entries, 0 to 1167

Data columns (total 80 columns):

#	Column	Non-Null Count	Dtype
0	MSSubClass	1168 non-null	int64
1	MSZoning	1168 non-null	object
2	LotFrontage	954 non-null	float64
3	LotArea	1168 non-null	int64
4	Street	1168 non-null	object
5	Alley	77 non-null	object
6	LotShape	1168 non-null	object
7	LandContour	1168 non-null	object
8	Utilities	1168 non-null	object
9	LotConfig	1168 non-null	object
10	LandSlope	1168 non-null	object
11	Neighborhood	1168 non-null	object
12	Condition1	1168 non-null	object
13	Condition2	1168 non-null	object
14	BldgType	1168 non-null	object
15	HouseStyle	1168 non-null	object
16	OverallQual	1168 non-null	int64
17	OverallCond	1168 non-null	int64
18	YearBuilt	1168 non-null	int64
19	YearRemodAdd	1168 non-null	int64
20	RoofStyle	1168 non-null	object
21	RoofMatl	1168 non-null	object
22	Exterior1st	1168 non-null	object
23	Exterior2nd	1168 non-null	object
24	MasVnrType	1161 non-null	object
25	MasVnrArea	1161 non-null	float64
26	ExterQual	1168 non-null	object
27	ExterCond	1168 non-null	object
28	Foundation	1168 non-null	object

29	BsmtQual	1138	non-null	object
30	BsmtCond	1138	non-null	object
31	BsmtExposure	1137	non-null	object
32	BsmtFinType1	1138	non-null	object
33	BsmtFinSF1	1168	non-null	int64
34	BsmtFinType2	1137	non-null	object
35	BsmtFinSF2	1168	non-null	int64
36	BsmtUnfSF	1168	non-null	int64
37	TotalBsmtSF	1168	non-null	int64
38	Heating	1168	non-null	object
39	HeatingQC	1168	non-null	object
40	CentralAir	1168	non-null	object
41	Electrical	1168	non-null	object
42	1stFlrSF	1168	non-null	int64
43	2ndFlrSF	1168	non-null	int64
44	LowQualFinSF	1168	non-null	int64
45	GrLivArea	1168	non-null	int64
46	BsmtFullBath	1168	non-null	int64
47	BsmtHalfBath	1168	non-null	int64
48	FullBath	1168	non-null	int64
49	HalfBath	1168	non-null	int64
50	BedroomAbvGr	1168	non-null	int64
51	KitchenAbvGr	1168	non-null	int64
52	KitchenQual	1168	non-null	object
53	TotRmsAbvGrd	1168	non-null	int64
54	Functional	1168	non-null	object
55	Fireplaces	1168	non-null	int64
56	FireplaceQu	617	non-null	object
57	GarageType	1104	non-null	object
58	GarageYrBlt	1104	non-null	float64
59	GarageFinish	1104	non-null	object
60	GarageCars	1168	non-null	int64
61	GarageArea	1168	non-null	int64
62	GarageQual	1104	non-null	object
63	GarageCond	1104	non-null	object
64	PavedDrive	1168	non-null	object
65	WoodDeckSF	1168	non-null	int64
66	OpenPorchSF	1168	non-null	int64
67	EnclosedPorch	1168	non-null	int64
68	3SsnPorch	1168	non-null	int64
69	ScreenPorch	1168	non-null	int64
70	PoolArea	1168	non-null	int64
71	PoolQC	7	non-null	object
72	Fence	237	non-null	object
73	MiscFeature	44	non-null	object
74	MiscVal	1168	non-null	int64
75	MoSold	1168	non-null	int64

```
76 YrSold      1168 non-null  int64
77 SaleType    1168 non-null  object
78 SaleCondition 1168 non-null  object
79 SalePrice   1168 non-null  int64
dtypes: float64(3), int64(34), object(43)
memory usage: 730.1+ KB
None
```

Check the Data type

```
print(test.dtypes)
print('*****')
print(train.dtypes)
```

```
MSSubClass      int64
MSZoning         object
LotFrontage     float64
LotArea          int64
Street          object
...
MiscVal          int64
MoSold           int64
YrSold           int64
SaleType         object
SaleCondition    object
Length: 79, dtype: object
*****
```

```
MSSubClass      int64
MSZoning         object
LotFrontage     float64
LotArea          int64
Street          object
...
MoSold           int64
YrSold           int64
SaleType         object
SaleCondition    object
SalePrice        int64
Length: 80, dtype: object
```

```
In [104]: print(test.dtypes)
          print('*****')
          print(train.dtypes)

MSSubClass      int64
MSZoning         object
LotFrontage     float64
LotArea          int64
Street          object
...
MiscVal          int64
MoSold           int64
YrSold           int64
SaleType         object
SaleCondition    object
Length: 79, dtype: object
*****
MSSubClass      int64
MSZoning         object
LotFrontage     float64
LotArea          int64
Street          object
...
MoSold           int64
YrSold           int64
SaleType         object
SaleCondition    object
SalePrice        int64
Length: 80, dtype: object
```

Observation: There are two types of data present in the dataset categorical and numerical.

Data Preprocessing Done

- I checked the information, data types, null values, correlation of the independent and dependent features and **from the correlation table**.
- Some columns can't have any negative value, so those columns were treated accordingly.
- Treated Null values accordingly columns type.
- Skewness, Outliers are treated manually for the features giving some important information, and then the threshold values were set to make the data free from outliers.
- Applied StandardScaler.
- Applied various machine learning model and compared it.

Handling Missing Values

Let's check the missing values of top 30 columns

```
print(train.isnull().values.any())  
print("*****")  
print(train.isnull().sum().sort_values(ascending = False).head(30))  
print("*****")  
print(test.isnull().sum().sort_values(ascending = False).head(30))
```

```
True  
*****  
PoolQC      1161  
MiscFeature  1124  
Alley       1091  
Fence       931  
FireplaceQu  551  
LotFrontage  214  
GarageType   64  
GarageFinish 64  
GarageQual   64  
GarageCond   64  
GarageYrBlt  64  
BsmtExposure 31  
BsmtFinType2 31  
BsmtCond     30  
BsmtFinType1 30  
BsmtQual     30  
MasVnrArea   7  
MasVnrType   7  
RoofStyle    0  
RoofMatl     0  
ExterQual    0  
Exterior1st  0  
Exterior2nd  0  
YearBuilt    0  
ExterCond    0  
Foundation   0  
YearRemodAdd 0
```

```

SalePrice      0
OverallCond    0
OverallQual    0
dtype: int64
*****
PoolQC         292
MiscFeature    282
Alley          278
Fence          248
FireplaceQu    139
LotFrontage    45
GarageCond     17
GarageType     17
GarageYrBlt    17
GarageFinish   17
GarageQual     17
BsmtFinType1   7
BsmtExposure   7
BsmtCond       7
BsmtQual       7
BsmtFinType2   7
Electrical     1
MasVnrArea     1
MasVnrType     1
LandSlope      0
RoofMatl       0
MSZoning       0
LotArea        0
Street         0
LotShape       0
Foundation     0
ExterCond      0
ExterQual      0
Exterior2nd    0
Exterior1st    0
dtype: int64

```

Observation:

In train dataset There are 18 columns that have missing values. Major missing values columns are PoolQC, 1124-in MiscFeature, 11091-in Alley, 931-in Fence, 551-in FireplaceQu

There are 1161-missing values in the column PoolQC, 1124-in MiscFeature, 11091-in Alley, 931-in Fence, 551-in FireplaceQu, 214- in LotFrontage, 64-each in GarageType, GarageCond, GarageYrBlt, GarageFinish, GarageQual, 31-in BsmtExposure and BsmtFinType2, 30-in BsmtCond and BsmtQual, 7-in MasVnrArea and MasVnrType present in our dataset.

IN test Dataset There are 19 columns that have missing values. Major missing values columns are PoolQC 292 MiscFeature 282 Alley 278 Fence 248 FireplaceQu 139

Total Missing Value Percantage for Train dataset

Your selected dataframe has 80 columns.
There are 18 columns that have missing values.

Out[107]:

	Missing Values	% of Total Values
PoolQC	1161	99.4
MiscFeature	1124	96.2
Alley	1091	93.4
Fence	931	79.7
FireplaceQu	551	47.2
LotFrontage	214	18.3
GarageType	64	5.5
GarageYrBlt	64	5.5
GarageFinish	64	5.5
GarageQual	64	5.5
GarageCond	64	5.5
BsmtExposure	31	2.7
BsmtFinType2	31	2.7
BsmtCond	30	2.6
BsmtFinType1	30	2.6
BsmtQual	30	2.6
MasVnrArea	7	0.6
MasVnrType	7	0.6

Total Missing Value Percentage for Test Dataset

t[108]:

	Missing Values	% of Total Values
PoolQC	292	100.0
MiscFeature	282	96.6
Alley	278	95.2
Fence	248	84.9
FireplaceQu	139	47.6
LotFrontage	45	15.4
GarageType	17	5.8
GarageYrBlt	17	5.8
GarageFinish	17	5.8
GarageQual	17	5.8
GarageCond	17	5.8
BsmtExposure	7	2.4
BsmtFinType1	7	2.4
BsmtFinType2	7	2.4
BsmtCond	7	2.4
BsmtQual	7	2.4
MasVnrArea	1	0.3
MasVnrType	1	0.3
Electrical	1	0.3

```
In [105]: # Let's explore the categorical columns

for column in train.columns:
    if train[column].dtypes == object:
        print(str(column) + ' : ' + str(train[column].unique()))
        print(train[column].value_counts())
        print('\n')
```

MSZoning : ['RL' 'RM' 'FV' 'RH' 'C (all)']

RL 928
RM 163
FV 52
RH 16
C (all) 9
Name: MSZoning, dtype: int64

Street : ['Pave' 'Grvl']
Pave 1164
Grvl 4
Name: Street, dtype: int64

Alley : [nan 'Grvl' 'Pave']
Grvl 41
Pave 36
Name: Alley, dtype: int64

LotShape : ['IR1' 'Reg' 'IR2' 'IR3']
Reg 740
IR1 390
IR2 32
IR3 6
Name: LotShape, dtype: int64

LandContour : ['Lvl' 'Bnk' 'HLS' 'Low']
Lvl 1046
Bnk 50
HLS 42
Low 30
Name: LandContour, dtype: int64

Utilities : ['AllPub']
AllPub 1168
Name: Utilities, dtype: int64

LotConfig : ['Inside' 'CulDSac' 'FR2' 'Corner' 'FR3']
Inside 842
Corner 222
CulDSac 69
FR2 33
FR3 2
Name: LotConfig, dtype: int64

LandSlope : ['Gtl' 'Mod' 'Sev']
Gtl 1105
Mod 51
Sev 12
Name: LandSlope, dtype: int64

Neighborhood : ['NPkVill' 'NAMES' 'NoRidge' 'NWAmes' 'Gilbert' 'Sawyer' 'Edwards'
'IDOTRR' 'CollgCr' 'Mitchel' 'Crawfor' 'BrDale' 'StoneBr' 'BrkSide'
'NridgHt' 'OldTown' 'Somerst' 'Timber' 'SWISU' 'SawyerW' 'ClearCr'
'Veenker' 'Blmngtn' 'MeadowV' 'Blueste']

NAMES	182
CollgCr	118
OldTown	86
Edwards	83
Somerst	68
Gilbert	64
NridgHt	61
Sawyer	60
NWAmes	59
SawyerW	51
BrkSide	50
Crawfor	45
NoRidge	35
Mitchel	34
IDOTRR	30
Timber	24
ClearCr	24
SWISU	21
StoneBr	19
Blmngtn	15
BrDale	11
MeadowV	9
Veenker	9
NPkVill	8
Blueste	2

Name: Neighborhood, dtype: int64

Condition1 : ['Norm' 'Feedr' 'RRAn' 'PosA' 'RRAe' 'Artery' 'PosN' 'RRNe' 'RRNn']

Norm	1005
Feedr	67
Artery	38
RRAn	20
PosN	17
RRAe	9
PosA	6
RRNn	4
RRNe	2

Name: Condition1, dtype: int64

Condition2 : ['Norm' 'RRAe' 'Feedr' 'PosN' 'Artery' 'RRNn' 'PosA' 'RRAn']

Norm	1154
Feedr	6
Artery	2
PosN	2
RRAe	1
RRNn	1
RRAn	1
PosA	1

Name: Condition2, dtype: int64

BldgType : ['TwnhsE' '1Fam' 'Duplex' 'Twnhs' '2fmCon']

1Fam 981

TwnhsE 90

Duplex 41

Twnhs 29

2fmCon 27

Name: BldgType, dtype: int64

HouseStyle : ['1Story' '2Story' '1.5Fin' 'SFoyer' '1.5Unf' 'SLvl' '2.5Fin' '2.5Unf']

1Story 578

2Story 361

1.5Fin 121

SLvl 47

SFoyer 32

1.5Unf 12

2.5Unf 10

2.5Fin 7

Name: HouseStyle, dtype: int64

RoofStyle : ['Gable' 'Flat' 'Hip' 'Shed' 'Gambrel' 'Mansard']

Gable 915

Hip 225

Flat 12

Gambrel 9

Mansard 5

Shed 2

Name: RoofStyle, dtype: int64

RoofMatl : ['CompShg' 'Tar&Grv' 'WdShngl' 'WdShake' 'Roll' 'ClyTile' 'Metal'
'Membran']

CompShg 1144

Tar&Grv 10

WdShngl 6

WdShake 4

Membran 1

Metal 1

ClyTile 1

Roll 1

Name: RoofMatl, dtype: int64

Exterior1st : ['Plywood' 'Wd Sdng' 'MetalSd' 'CemntBd' 'VinylSd' 'HdBoard' 'Stucco'
'WdShing' 'BrkFace' 'Stone' 'AsbShng' 'AsphShn' 'ImStucc' 'BrkComm']

VinylSd 396

HdBoard 179

MetalSd 178

Wd Sdng 174

Plywood 93

CemntBd 42

BrkFace 41

Stucco 22

WdShing 19

AsbShng 19

Stone 2

AsphShn 1
BrkComm 1
ImStucc 1
Name: Exterior1st, dtype: int64

Exterior2nd : ['Plywood' 'Wd Sdng' 'MetalSd' 'CmentBd' 'VinylSd' 'HdBoard' 'Wd Shng'
'Stucco' 'ImStucc' 'Stone' 'BrkFace' 'AsbShng' 'Brk Cmn' 'AsphShn'
'Other']
VinylSd 387
MetalSd 173
HdBoard 170
Wd Sdng 165
Plywood 118
CmentBd 42
Wd Shng 31
Stucco 23
BrkFace 20
AsbShng 18
ImStucc 8
Brk Cmn 5
Stone 4
AsphShn 3
Other 1
Name: Exterior2nd, dtype: int64

MasVnrType : ['None' 'BrkFace' 'Stone' 'BrkCmn' nan]
None 696
BrkFace 354
Stone 98
BrkCmn 13
Name: MasVnrType, dtype: int64

ExterQual : ['TA' 'Gd' 'Ex' 'Fa']
TA 717
Gd 397
Ex 43
Fa 11
Name: ExterQual, dtype: int64

ExterCond : ['TA' 'Gd' 'Fa' 'Po' 'Ex']
TA 1022
Gd 117
Fa 26
Ex 2
Po 1
Name: ExterCond, dtype: int64

Foundation : ['CBlock' 'PConc' 'BrkTil' 'Slab' 'Stone' 'Wood']
CBlock 516
PConc 513
BrkTil 112
Slab 21

Stone 5
Wood 1
Name: Foundation, dtype: int64

BsmtQual : ['Gd' 'TA' 'Ex' nan 'Fa']
TA 517
Gd 498
Ex 94
Fa 29
Name: BsmtQual, dtype: int64

BsmtCond : ['TA' 'Gd' 'Fa' nan 'Po']
TA 1041
Gd 56
Fa 39
Po 2
Name: BsmtCond, dtype: int64

BsmtExposure : ['No' 'Gd' 'Av' 'Mn' nan]
No 756
Av 180
Gd 108
Mn 93
Name: BsmtExposure, dtype: int64

BsmtFinType1 : ['ALQ' 'GLQ' 'BLQ' 'Unf' 'Rec' 'LwQ' nan]
Unf 345
GLQ 330
ALQ 174
BLQ 121
Rec 109
LwQ 59
Name: BsmtFinType1, dtype: int64

BsmtFinType2 : ['Unf' 'Rec' 'BLQ' 'GLQ' nan 'ALQ' 'LwQ']
Unf 1002
Rec 43
LwQ 40
BLQ 24
ALQ 16
GLQ 12
Name: BsmtFinType2, dtype: int64

Heating : ['GasA' 'GasW' 'Floor' 'OthW' 'Wall' 'Grav']
GasA 1143
GasW 14
Grav 5
Wall 4
Floor 1
OthW 1
Name: Heating, dtype: int64

HeatingQC : ['TA' 'Ex' 'Gd' 'Fa' 'Po']

Ex 585

TA 352

Gd 192

Fa 38

Po 1

Name: HeatingQC, dtype: int64

CentralAir : ['Y' 'N']

Y 1090

N 78

Name: CentralAir, dtype: int64

Electrical : ['SBrkr' 'FuseA' 'FuseF' 'FuseP' 'Mix']

SBrkr 1070

FuseA 74

FuseF 21

FuseP 2

Mix 1

Name: Electrical, dtype: int64

KitchenQual : ['TA' 'Gd' 'Ex' 'Fa']

TA 578

Gd 478

Ex 82

Fa 30

Name: KitchenQual, dtype: int64

Functional : ['Typ' 'Mod' 'Maj1' 'Min1' 'Min2' 'Sev' 'Maj2']

Typ 1085

Min2 30

Min1 25

Mod 12

Maj1 11

Maj2 4

Sev 1

Name: Functional, dtype: int64

FireplaceQu : ['TA' 'Gd' nan 'Fa' 'Ex' 'Po']

Gd 301

TA 252

Fa 25

Ex 21

Po 18

Name: FireplaceQu, dtype: int64

GarageType : ['Attchd' 'BuiltIn' 'Detchd' 'Basment' nan '2Types' 'CarPort']

Attchd 691

Detchd 314

BuiltIn 70
Basment 16
CarPort 8
2Types 5
Name: GarageType, dtype: int64

GarageFinish : ['RFn' 'Unf' 'Fin' nan]
Unf 487
RFn 339
Fin 278
Name: GarageFinish, dtype: int64

GarageQual : ['TA' 'Fa' nan 'Gd' 'Ex' 'Po']
TA 1050
Fa 39
Gd 11
Po 2
Ex 2
Name: GarageQual, dtype: int64

GarageCond : ['TA' 'Fa' 'Gd' nan 'Po' 'Ex']
TA 1061
Fa 28
Gd 8
Po 6
Ex 1
Name: GarageCond, dtype: int64

PavedDrive : ['Y' 'N' 'P']
Y 1071
N 74
P 23
Name: PavedDrive, dtype: int64

PoolQC : [nan 'Ex' 'Gd' 'Fa']
Gd 3
Fa 2
Ex 2
Name: PoolQC, dtype: int64

Fence : [nan 'MnPrv' 'GdPrv' 'GdWo' 'MnWw']
MnPrv 129
GdPrv 51
GdWo 47
MnWw 10
Name: Fence, dtype: int64

MiscFeature : [nan 'Shed' 'Gar2' 'TenC' 'Othr']
Shed 40
Gar2 2

TenC 1
Othr 1
Name: MiscFeature, dtype: int64

SaleType : ['WD' 'COD' 'New' 'ConLI' 'ConLw' 'Con' 'ConLD' 'Oth' 'CWD']
WD 999
New 106
COD 38
ConLD 8
ConLI 5
ConLw 4
Oth 3
CWD 3
Con 2
Name: SaleType, dtype: int64

SaleCondition : ['Normal' 'Partial' 'Abnorml' 'Family' 'Alloca' 'AdjLand']
Normal 945
Partial 108
Abnorml 81
Family 18
Alloca 12
AdjLand 4
Name: SaleCondition, dtype: int64

Observation:

There is only one unique value present in utilities column so we will be dropping this column.

2. In categorical columns there are missing values present in columns Alley, MasVnrType, BsmtQual, BsmtCond, BsmtExposure, BsmtFinType1, BsmtFinType2, FireplaceQu, GarageType, GarageFinish, GarageQual, GarageCond, PoolQC, Fence, MiscFeature

Fill Missing Values

```
In [111]: # Let's fill the missing values in categorical columns as NA in train dataset

columns = ["FireplaceQu", "GarageType", "GarageFinish", "GarageQual", "GarageCond", "BsmtExposure", "BsmtFinType2", "BsmtCond", "
train[columns] = train[columns].fillna('NA')

# Let's fill the missing values in MasVnrType with None
train['MasVnrType'] = train['MasVnrType'].fillna('None')

# Let's fill the missing values in GarageYrBlt with 0
train['GarageYrBlt'] = train['GarageYrBlt'].fillna('0')

# Let's Imputing the missing values and replace it with the median
train['LotFrontage'].fillna(train['LotFrontage'].median(),inplace=True)
train['MasVnrArea'].fillna(train['MasVnrArea'].median(),inplace=True)

In [112]: # Let's fill the missing values in categorical columns as NA in test dataset

columns = ["FireplaceQu", "GarageType", "GarageFinish", "GarageQual", "GarageCond", "BsmtExposure", "BsmtFinType2", "BsmtCond", "
test[columns] = test[columns].fillna('NA')

# Let's fill the missing values in MasVnrType with None
test['MasVnrType'] = test['MasVnrType'].fillna('None')

# Let's fill the missing values in GarageYrBlt with 0
test['GarageYrBlt'] = test['GarageYrBlt'].fillna('0')

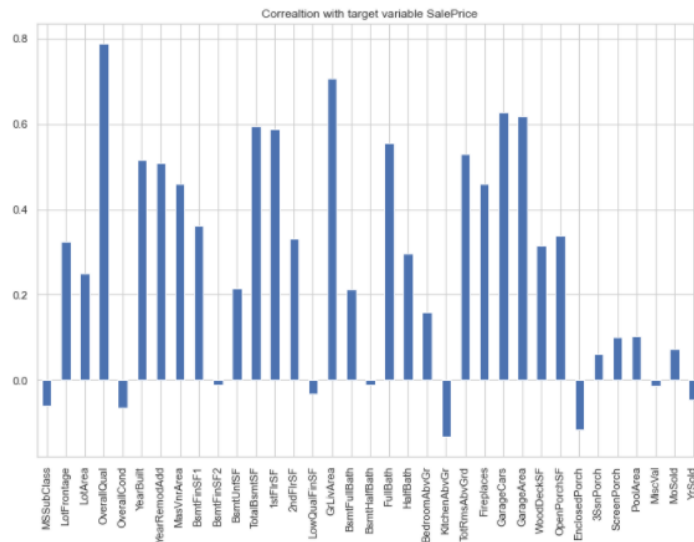
# Let's Imputing the missing values and replace it with the median
test['LotFrontage'].fillna(test['LotFrontage'].median(),inplace=True)
test['MasVnrArea'].fillna(test['MasVnrArea'].median(),inplace=True)
```

Data Inputs- Logic- Output Relationships

Lets check the correlation with target variable "Salesprice".

```
In [118]: # Let's check the correlation with target variable 'SalePrice'

plt.figure(figsize=(12,8))
train.drop('SalePrice', axis=1).corrwith(train['SalePrice']).plot(kind='bar',grid=True)
plt.xticks(rotation='vertical')
plt.title("Correaltion with target variable SalePrice");
```



Observation:

Observation:

1. The column OverallQual is most positively correlated with SalePrice.
- 2.The column KitchenAbvGrd and EnclosedPorch is most negatively correlated with SalePrice.

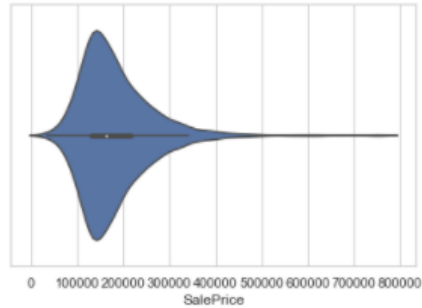
Maximum number of SalePrice lies between 140000 and 230000.

Univariate Analysis

In [119]: *# Let's Check the target variable*

```
sns.set(style='whitegrid')
sns.violinplot(train['SalePrice'])
plt.show()

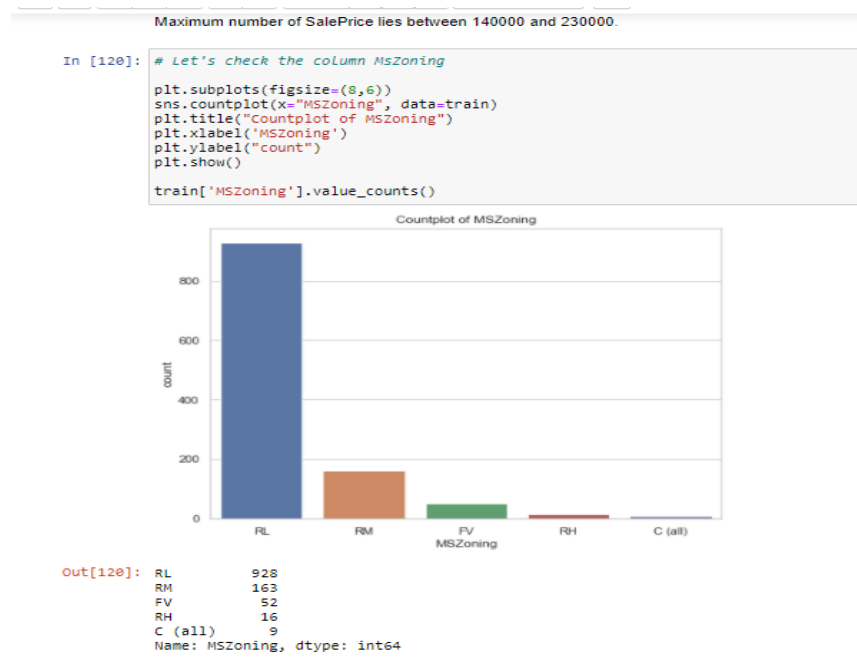
train['SalePrice'].value_counts()
```



```
Out[119]: 140000    18
          135000    16
          155000    12
          139000    11
          160000    11
          ..
          126175     1
          204000     1
          186000     1
          369900     1
          105500     1
          Name: SalePrice, Length: 581, dtype: int64
```

Maximum number of SalePrice lies between 140000 and 230000.

Maximum, 928 number of MSZoning are RL



Bivariate Analysis

Let's plot the Scatter plot between all feature variables and target variable

for col in train.describe().columns:

```
data=train.copy()
```

```
plt.scatter(data[col],data['SalePrice'])
```

```
plt.xlabel(col)
```

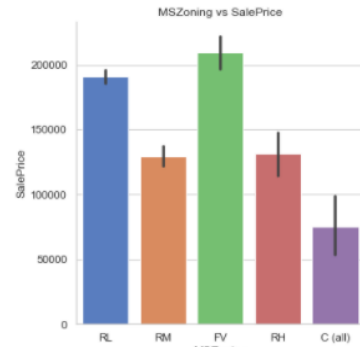
```
plt.ylabel('SalePrice')
```

```
plt.show()
```


1. SalePrice is maximum with FV MSZoning.

```
In [129]: # Let's plot the Factor plot of MSZoning vs SalePrice
plt.figure(figsize=(8,6))
sns.factorplot(x='MSZoning',y='SalePrice',data=train,kind='bar',size=5,palette='muted',aspect=1)
plt.title('MSZoning vs SalePrice')
plt.ylabel('SalePrice')
plt.show()
print(train.groupby('SalePrice')['MSZoning'].value_counts());
```

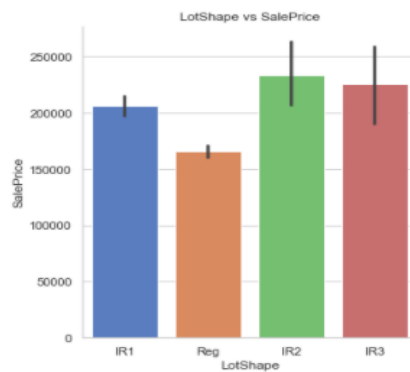
<Figure size 576x432 with 0 Axes>



2. SalePrice is maximum with IR2 LotShape.

```
In [130]: # Let's plot the Factor plot of LotShape vs SalePrice
plt.figure(figsize=(8,6))
sns.factorplot(x='LotShape',y='SalePrice',data=train,kind='bar',size=5,palette='muted',aspect=1)
plt.title('LotShape vs SalePrice')
plt.ylabel('SalePrice')
plt.show()
print(train.groupby('SalePrice')['LotShape'].value_counts());
```

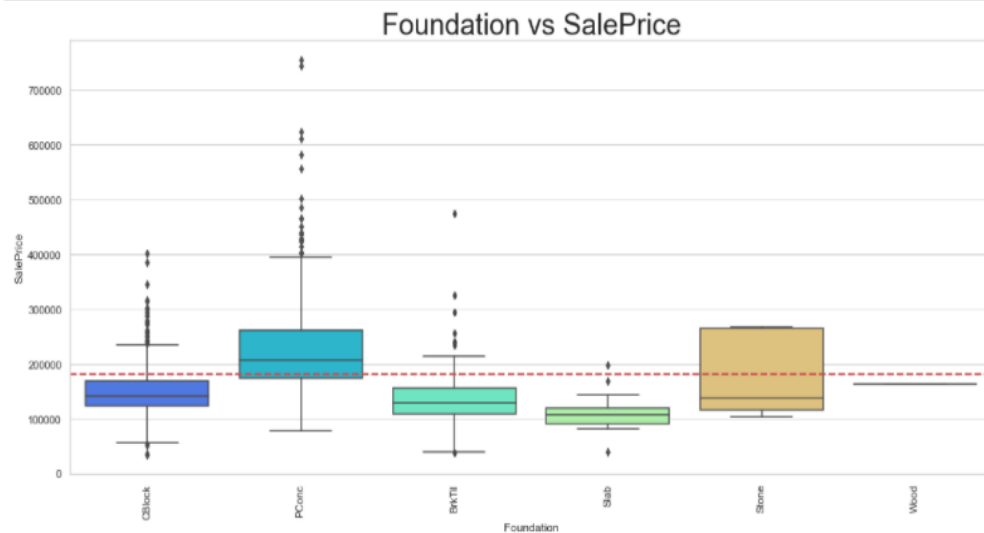
<Figure size 576x432 with 0 Axes>



3. SalePrice is maximum with PConc

```
In [133]: # Let's plot the Foundation vs SalePrice plot
```

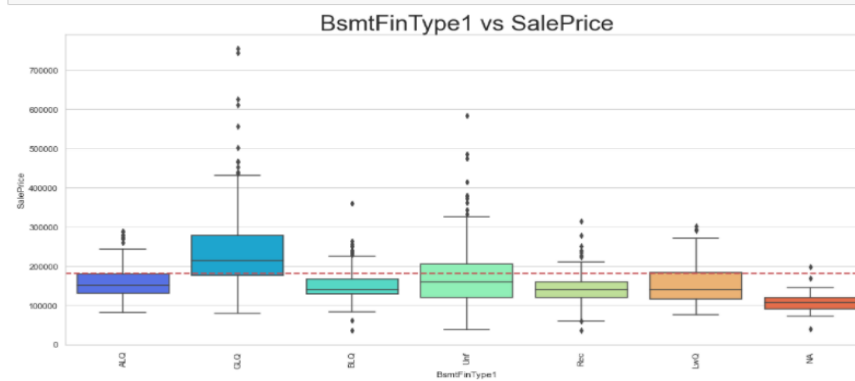
```
plt.figure(figsize=(18,8))
mean_price=np.mean(train['SalePrice'])
sns.boxplot(y='SalePrice',x='Foundation',data=train,palette="rainbow")
plt.axhline(mean_price,color='r',linestyle='dashed',linewidth=2)
plt.title("Foundation vs SalePrice",fontsize=30)
plt.xticks(rotation='vertical')
plt.show()
```



4. SalePrice is maximum with GLQ BsmtFinType1.

```
In [134]: # Let's plot the BsmtFinType1 vs SalePrice plot
```

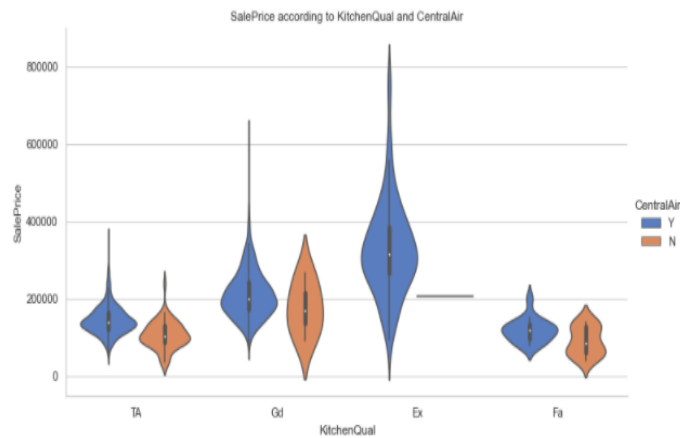
```
plt.figure(figsize=(18,8))
mean_price=np.mean(train['SalePrice'])
sns.boxplot(y='SalePrice',x='BsmtFinType1',data=train,palette="rainbow")
plt.axhline(mean_price,color='r',linestyle='dashed',linewidth=2)
plt.title("BsmtFinType1 vs SalePrice",fontsize=30)
plt.xticks(rotation='vertical')
plt.show()
```



Multivariate Analysis

```
In [135]: # Let's plot the GarageType and GarageCond with respect to SalePrice plot
```

```
sns.factorplot(x='KitchenQual', y='SalePrice', hue='CentralAir', data=train, kind='violin', size=5, palette='muted', aspect=2)
plt.title('SalePrice according to KitchenQual and CentralAir')
plt.xticks()
plt.ylabel('SalePrice')
plt.show()
```

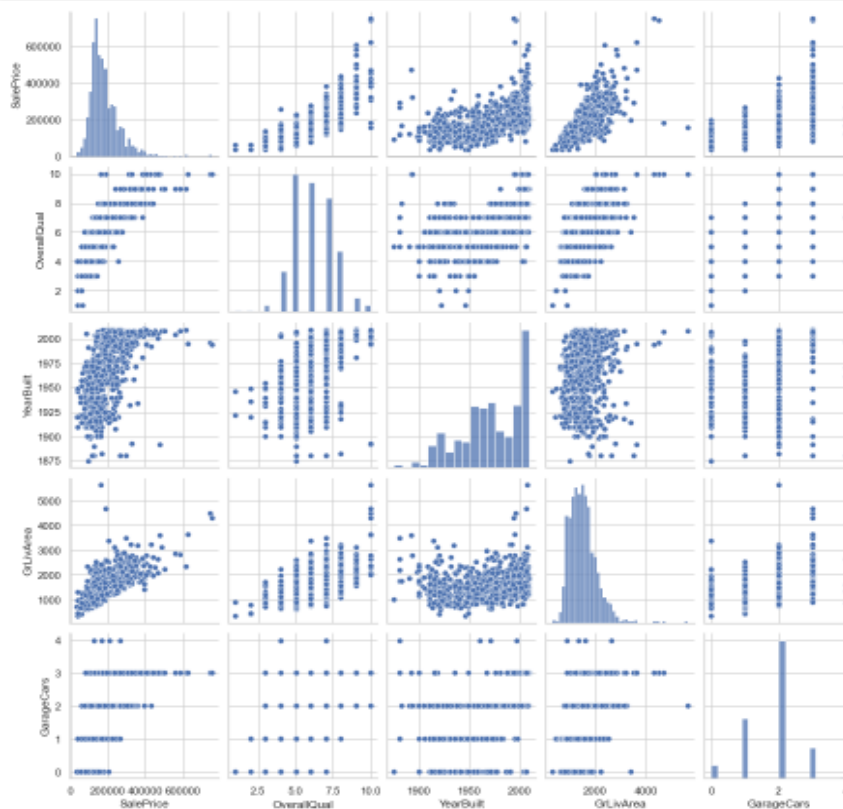


Let's plot the pairplot

```
sns.pairplot(train,
vars=['SalePrice','OverallQual','YearBuilt','GrLivArea','GarageCars']);
```

```
In [136]: # Let's plot the pairplot
```

```
sns.pairplot(train, vars=["SalePrice", "OverallQual", "YearBuilt", "GrLivArea", "GarageCars"]);
```



- State the set of assumptions (if any) related to the problem under consideration
I have not consider any pre-assumption , project performance from beginning to end is based on data facts only.

• Hardware and Software Requirements and Tools Used

Windows Edition-Windows 8.1 Pro
Processor-Intel(R) Core(TM) i3-5005U CPU @ 2.00GHz 2.00GHz
Installed memory RAM- 4 GB
System Type-64 bit OS, x64 based processor

Software Requirement- Anaconda 4.9.2 , Python 3.8.5, Jupiter Notebook.

Libraries used:-

```
In [ ]: # Let's import all the required Libraries

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
pd.pandas.set_option('display.max_columns',None)

from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from scipy import stats

from sklearn.metrics import mean_absolute_error
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
from sklearn import linear_model
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split

from sklearn.linear_model import LinearRegression,Lasso,Ridge,ElasticNet
from sklearn.svm import SVR
from sklearn.neighbors import KNeighborsRegressor
from sklearn.tree import DecisionTreeRegressor

from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import AdaBoostRegressor
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.model_selection import GridSearchCV,cross_val_score
from sklearn.model_selection import GridSearchCV

#importing warnings
import warnings
warnings.filterwarnings('ignore')
```

Model/s Development and Evaluation

- Identification of possible problem-solving approaches (methods)
- **Analytical Approach** –Based on type of data by performing EDA I have decided which model to be used for this data.
- **Statistical Approach** – Data should be in scaled manner, it should not be distorted, for that all values using mean method due to continuous data numbers.

Statistical Approach for train Dataset

```
In [113]: # Let's check the statistical summary of our dataset
train.describe()
```

Out[113]:

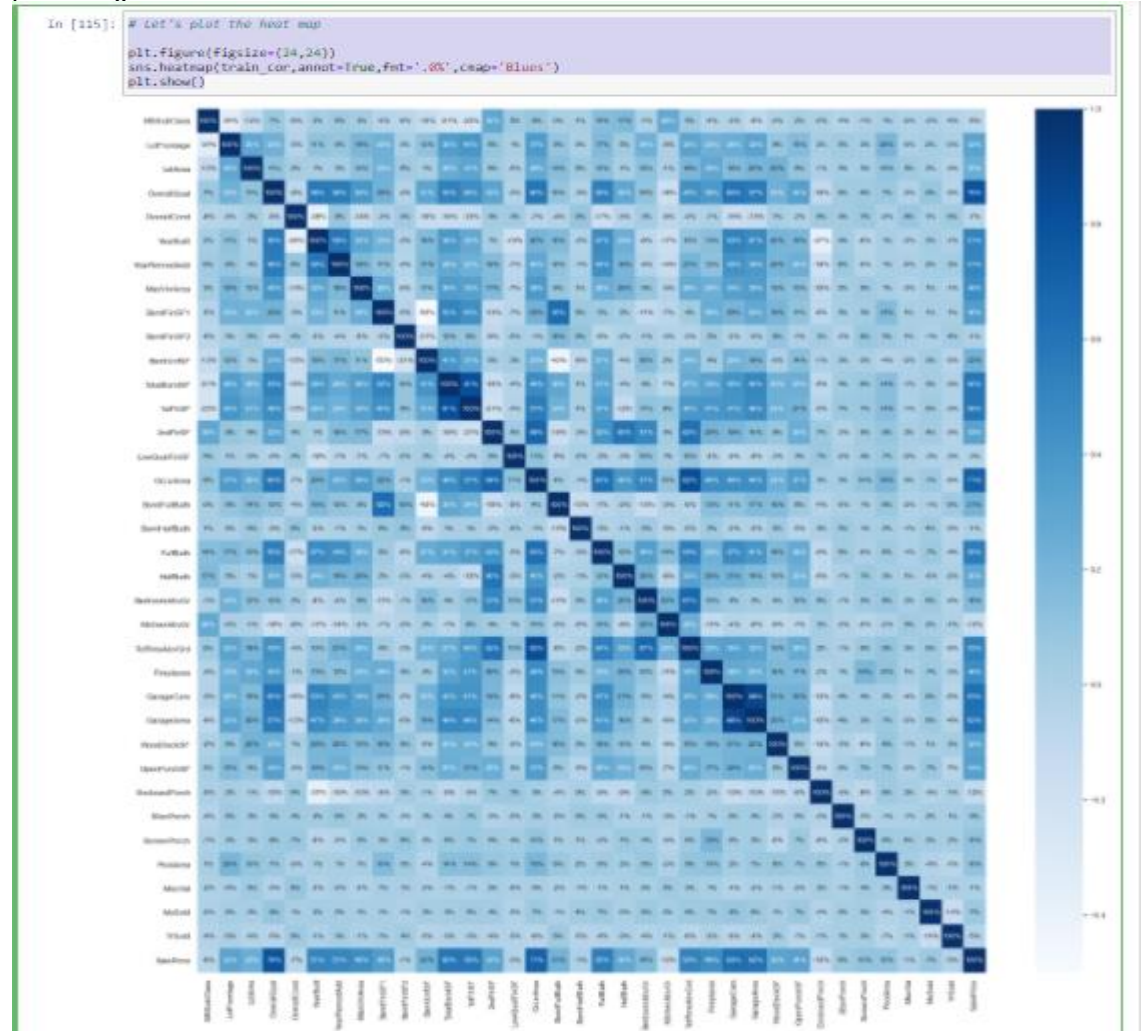
	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd	MasVnrArea	BsmtFinSF1	BsmtFinSF2	BsmtUnfSF
count	1168.000000	1168.000000	1168.000000	1168.000000	1168.000000	1168.000000	1168.000000	1168.000000	1168.000000	1168.000000	1168.000000
mean	56.767979	70.807363	10484.749144	6.104452	5.595890	1970.930651	1984.758562	101.696918	444.726027	46.647260	569.721747
std	41.940650	22.440317	8957.442311	1.390153	1.124343	30.145255	20.785185	182.218483	462.664785	163.520016	449.375525
min	20.000000	21.000000	1300.000000	1.000000	1.000000	1875.000000	1950.000000	0.000000	0.000000	0.000000	0.000000
25%	20.000000	60.000000	7621.500000	5.000000	5.000000	1954.000000	1966.000000	0.000000	0.000000	0.000000	216.000000
50%	50.000000	70.000000	9522.500000	6.000000	5.000000	1972.000000	1993.000000	0.000000	385.500000	0.000000	474.000000
75%	70.000000	79.250000	11515.500000	7.000000	6.000000	2000.000000	2004.000000	160.000000	714.500000	0.000000	816.000000
max	190.000000	313.000000	164660.000000	10.000000	9.000000	2010.000000	2010.000000	1600.000000	5644.000000	1474.000000	2336.000000

- Maximum standard deviation of 8957.44 is observed in LotArea column.
- Maximum SalePrice of a house observed is 755000 and minimum is 34900.
 - In the columns Id, MSSubclass, LotArea, MasVnrArea, BsmtFinSF1, BsmtFinSF2, BsmtUnfSF, TotalBsmtSF, 1stFlrSF, 2ndFlrSF, LowQualFinSF, GrLivArea, BsmtFullBath, HalfBath, TotRmsAbvGrd, WoodDeckSF, OpenPorchSF, EnclosedPorch, 3SsnPorch, ScreenPorch, PoolArea, Miscval, salePrice mean is considerably greater than median so the columns are positively skewed.
 - In the columns FullBath, BedroomAbvGr, Fireplaces, Garagecars, GarageArea, YrSold Median is greater than mean so the columns are negatively skewed.
 - In the columns Id, MSSubClass, LotFrontage, LotArea, MasVnrArea, BsmtFinSF1, BsmtFinSF2, BsmtUnfSF, TotalBsmtSF, 1stFlrSF, 2ndFlrSF, LowQualFinSF, GrLivArea, BsmtHalfBath, BedroomAbvGr, ToRmsAbvGrd, GarageArea, WoodDeckSF, OpenPorchSF, EnclosedPorch, 3SsnPorch, ScreenPorch, PoolArea, MiscVal, SalePrice there is considerable difference between the 75 percentile and maximum so outliers are present.

Checking Correlation with Heatmap

Let's plot the heat map

```
plt.figure(figsize=(24,24))
sns.heatmap(train_cor,annot=True,fmt='.0%',cmap='Blues')
plt.show()
```



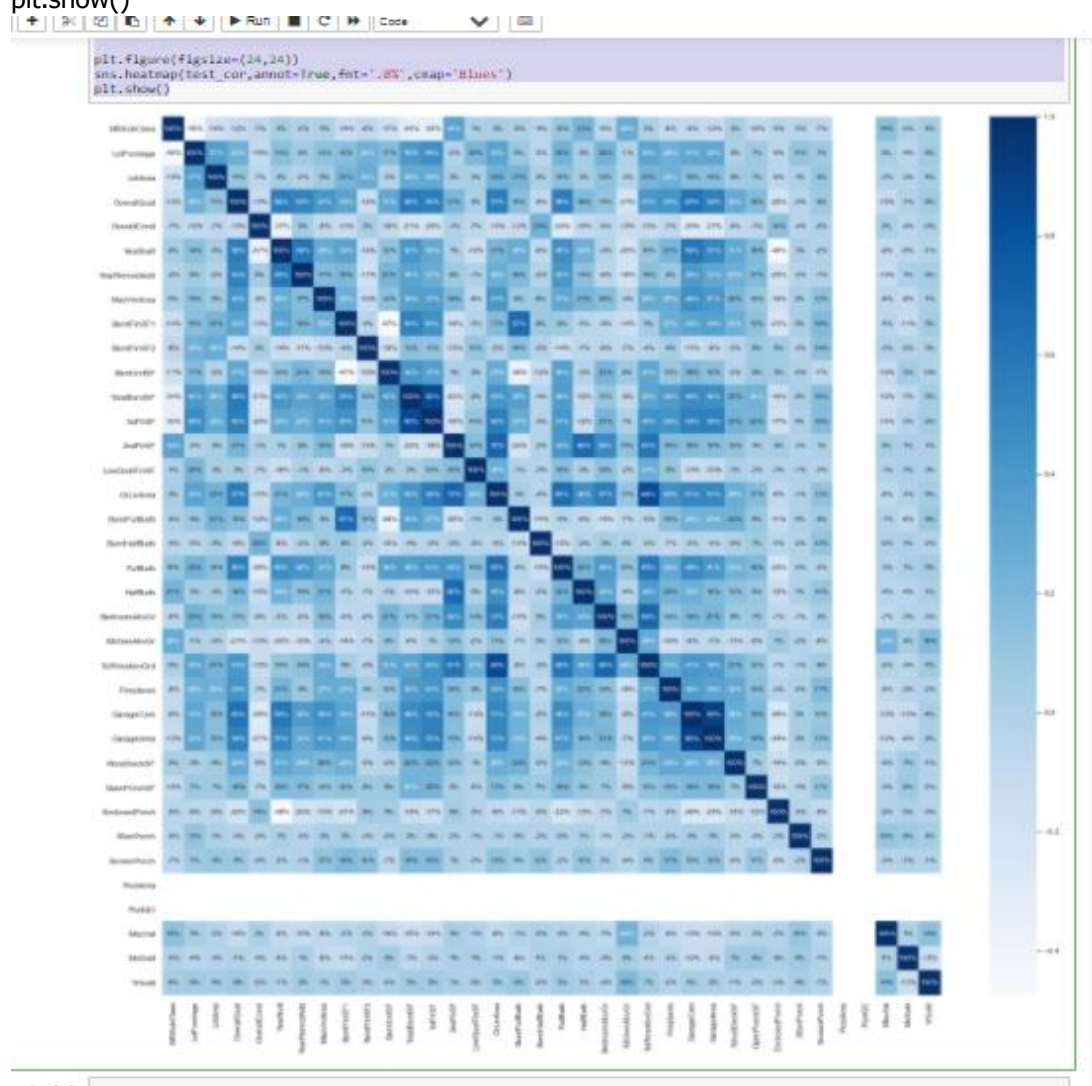
1.SalePrice is highly positively correlated with the columns OverallQual, YearBuilt, YearRemodAdd, TotalBsmtSF, 1stFlrSF, GrLivArea, FullBath, TotRmsAbvGrd, GarageCars, GarageArea.

2.SalePrice is negatively correlated with OverallCond, KitchenAbvGr, Encloseporch, YrSold.

3.We observe multicollinearity in between columns so we will be using Principal Component Analysis(PCA).

```
# Let's plot the heat map for test dataset
```

```
plt.figure(figsize=(24,24))
sns.heatmap(test_cor,annot=True,fmt='.0%',cmap='Blues')
plt.show()
```



Handling Outliers and skewness

Handling outliers and skewness

```
In [142]: # Let's make a copy of our dataset
train_cap = train.copy()

In [143]: def percentile_capping(train, cols, from_low_end, from_high_end):
    for col in cols:
        stats.mstats.winsorize(a=train[col], limits=(from_low_end, from_high_end), inplace=True)

In [144]: features=['MSSubClass', 'LotFrontage', 'LotArea', 'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd', 'MasVnrArea', 'BsmtFi

In [145]: percentile_capping(train, features, 0.01, 0.10)
# Let's check the shape
train_cap.shape

Out[145]: (1168, 244)

In [146]: for col in features:
    plt.figure(figsize=(16,4))

    plt.subplot(141)
    sns.distplot(train[col], label="skew: " + str(np.round(train[col].skew(),2)))
    plt.title('Before')
    plt.legend()

    plt.subplot(142)
    sns.distplot(train_cap[col], label="skew: " + str(np.round(train_cap[col].skew(),2)))
    plt.title('After')
    plt.legend()

    plt.subplot(143)
    sns.boxplot(train[col])
    plt.title('Before')

    plt.subplot(144)
    sns.boxplot(train_cap[col])
    plt.title('After')
    plt.tight_layout()
    plt.show()
```

• Testing of Identified Approaches (Algorithms)

All Algorithms list

```
from sklearn.metrics import mean_absolute_error
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
from sklearn import linear_model
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
```

```
from sklearn.linear_model import LinearRegression,Lasso,Ridge,ElasticNet
from sklearn.svm import SVR
from sklearn.neighbors import KNeighborsRegressor
from sklearn.tree import DecisionTreeRegressor
```

```
from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import AdaBoostRegressor
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.model_selection import GridSearchCV,cross_val_score
from sklearn.model_selection import GridSearchCV
```



```
from sklearn.metrics import mean_absolute_error
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
from sklearn import linear_model
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split

from sklearn.linear_model import LinearRegression,Lasso,Ridge,ElasticNet
from sklearn.svm import SVR
from sklearn.neighbors import KNeighborsRegressor
from sklearn.tree import DecisionTreeRegressor

from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import AdaBoostRegressor
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.model_selection import GridSearchCV,cross_val_score
from sklearn.model_selection import GridSearchCV
```

Below are Linear Regression algorithms used for the training and testing this dataset.

```
model=[LinearRegression(),
        DecisionTreeRegressor(),
        KNeighborsRegressor(),
        SVR(),
        Lasso(),
        Ridge(),
        ElasticNet(),
        RandomForestRegressor(),
        AdaBoostRegressor(),
```

```
In [167]: model=[LinearRegression(),
                  DecisionTreeRegressor(),
                  KNeighborsRegressor(),
                  SVR(),
                  Lasso(),
                  Ridge(),
                  ElasticNet(),
                  RandomForestRegressor(),
                  AdaBoostRegressor(),
```

a. Run and Evaluate selected models

PCA

```
In [154]: # Let's explore the PCA train dataset
```

```
covar_matrix = PCA(n_components = len(x.columns))  
covar_matrix.fit(x)
```

```
Out[154]: PCA(n_components=243)
```

```
In [155]: # Let's check the PCA test dataset
```

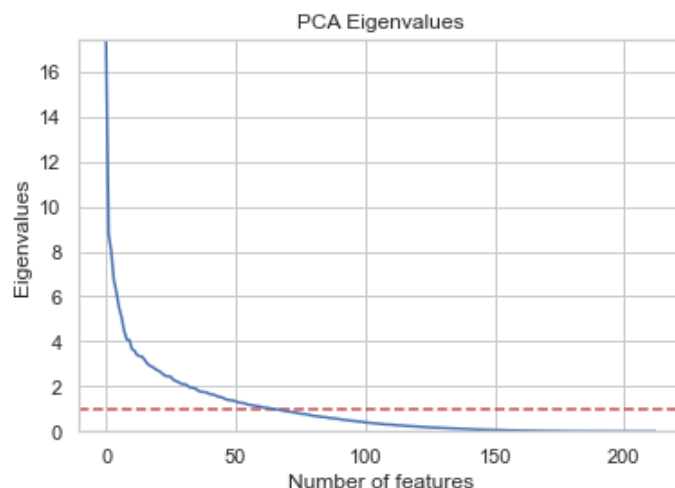
```
covar_matrix = PCA(n_components = len(x1.columns))  
covar_matrix.fit(x1)
```

```
Out[155]: PCA(n_components=213)
```

```
In [156]: # Let's plot the PCA componenets
```

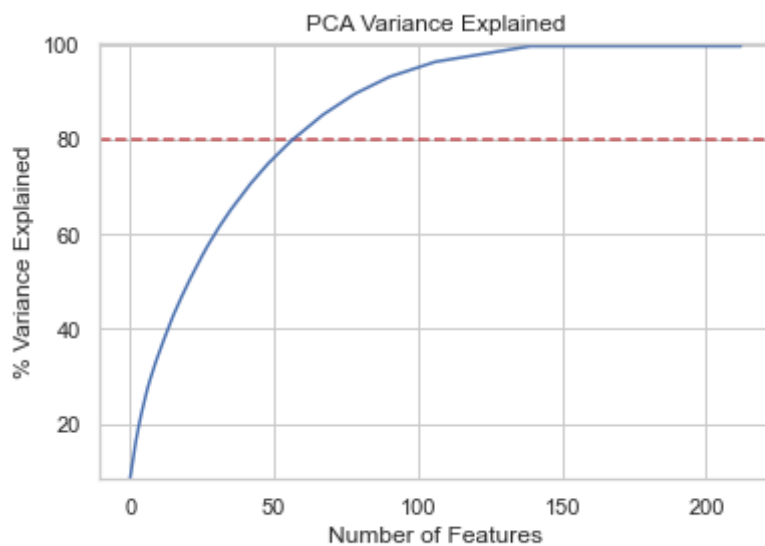
```
plt.ylabel('Eigenvalues')  
plt.xlabel('Number of features')  
plt.title('PCA Eigenvalues')  
plt.ylim(0,max(covar_matrix.explained_variance_))  
plt.style.context('seaborn-whitegrid')  
plt.axhline(y=1, color='r', linestyle='--')  
plt.plot(covar_matrix.explained_variance_)  
plt.show()
```

```
plt.show()
```



```
In [157]: variance = covar_matrix.explained_variance_ratio_
var=np.cumsum(np.round(covar_matrix.explained_variance_ratio_, decimals=3)*100)

plt.ylabel('% Variance Explained')
plt.xlabel('Number of Features')
plt.title('PCA Variance Explained')
plt.ylim(min(var),100.5)
plt.style.context('seaborn-whitegrid')
plt.axhline(y=80, color='r', linestyle='--')
plt.plot(var)
plt.show()
```



Lets find the best Random forest score

```
In [162]: # Let's find the best random state

max_r_score=0
for r_state in range(1,100):
    x_train,x_test,y_train,y_test=train_test_split(x,y,random_state=r_state,test_size=0.20)
    regr=linear_model.LinearRegression()
    regr.fit(x_train,y_train)
    y_pred=regr.predict(x_test)
    r2_scr=r2_score(y_test,y_pred)
    if r2_scr>max_r_score:
        max_r_score=r2_scr
        final_r_state=r_state
print("max r2 score corresponding to",final_r_state,"is",max_r_score)

max r2 score corresponding to 48 is 0.8496659416265843
```

max r2 score corresponding to 48 is 0.8496659416265843

```
In [163]: # Let's split the dataset into test and train
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.20,random_state=54)

In [167]: model=[LinearRegression(),
                  DecisionTreeRegressor(),
                  KNeighborsRegressor(),
                  SVR(),
                  Lasso(),
                  Ridge(),
                  ElasticNet(),
                  RandomForestRegressor(),
                  AdaBoostRegressor(),
                  ]
for m in model:
    m.fit(x_train,y_train)
    print('score of',m,'is:',m.score(x_train,y_train))
    predm=m.predict(x_test)
    print('Error:')
    print('Mean absolute error:',mean_absolute_error(y_test,predm))
    print('Mean squared error:',mean_squared_error(y_test,predm))
    print('Root Mean Squared Error:',np.sqrt(mean_squared_error(y_test,predm)))
    print("r2_score:",r2_score(y_test,predm))
    print('*****')
    print('\n')
```

score of LinearRegression() is: 0.82477807880984

Error:

Mean absolute error: 21214.668674560457

Mean squared error: 988711234.5151851

Root Mean Squared Error: 31443.778947753482

r2_score: 0.8493293024469674

score of DecisionTreeRegressor() is: 1.0

Error:

Mean absolute error: 32526.239316239316

Mean squared error: 2655740464.3076925

Root Mean Squared Error: 51533.87686083488

r2_score: 0.5952890446589649

score of KNeighborsRegressor() is: 0.8005580881741488

Error:

Mean absolute error: 26287.979487179484

Mean squared error: 1638191388.2099144

Root Mean Squared Error: 40474.57706029693

r2_score: 0.7503543698398546

score of SVR() is: -0.04568255380776742

Error:

Mean absolute error: 58256.27581850125

Mean squared error: 6883564961.965359

Root Mean Squared Error: 82967.2523467263

r2_score: -0.04899337467274023

score of Lasso() is: 0.8247780692387235

Error:

Mean absolute error: 21212.494199237746

Mean squared error: 988627101.9067537

Root Mean Squared Error: 31442.441093317702
r2_score: 0.8493421234996235

score of Ridge() is: 0.8247780154162465
Error:
Mean absolute error: 21207.71025561312
Mean squared error: 988429671.2769494
Root Mean Squared Error: 31439.301380230278
r2_score: 0.8493722101514916

score of ElasticNet() is: 0.8174701175087531
Error:
Mean absolute error: 19910.440060332676
Mean squared error: 1002174977.3116897
Root Mean Squared Error: 31657.147333764766
r2_score: 0.8472775491665279

score of RandomForestRegressor() is: 0.96680098429731
Error:
Mean absolute error: 21953.085384615388
Mean squared error: 1164063400.6093924
Root Mean Squared Error: 34118.37335819796
r2_score: 0.8226072098272706

score of AdaBoostRegressor() is: 0.8380210963982864
Error:
Mean absolute error: 30931.247989936797
Mean squared error: 1749721542.0135958
Root Mean Squared Error: 41829.67298477954
r2_score: 0.7333581777413218

Ridge is giving us minimum Rmse score so we choose it as our final model.

- **Key Metrics for success in solving problem under consideration**
- Key Metrics used were the Lasso, ridge, Elasticnet to find r2 Score and GridsearchCV score as this was Linear Regression problem and we focus more on R2score metrics to observe Mean absolute error, Mean squared error and Root Mean Squared Error.

- **Visualizations**

Hyperparameter tuning

```
In [79]: # Let's Use the GridSearchCV to find the best parameters in Ridge Regressor

parameters={'alpha': [25,10,4,2,1.0,0.8,0.5,0.3,0.2,0.1,0.05,0.02,0.01]}
rg=Ridge()

reg=GridSearchCV(rg,parameters,n_jobs=-1)
reg.fit(x,y)
print(reg.best_params_)

{'alpha': 25}
```

```
In [80]: # Let's use the Ridge Regressor with its best parameters

RG=Ridge(alpha=25)
RG.fit(x_train,y_train)
print('Score:',RG.score(x_train,y_train))
y_pred=RG.predict(x_test)
print('\n')
print('Mean absolute error:',mean_absolute_error(y_test,y_pred))
print('Mean squared error:',mean_squared_error(y_test,y_pred))
print('Root Mean Squared error:',np.sqrt(mean_squared_error(y_test,y_pred)))
print('\n')
print("r2_score:",r2_score(y_test,y_pred))
print('\n')

Score: 0.8325270027518725
```

Score: 0.8325270027518725

Mean absolute error: 20277.903007353238

Mean squared error: 921483021.3925847

Root Mean Squared error: 30355.938815865746

r2_score: 0.8595742773322715

```
In [81]: # Let's Cross validate the Ridge model

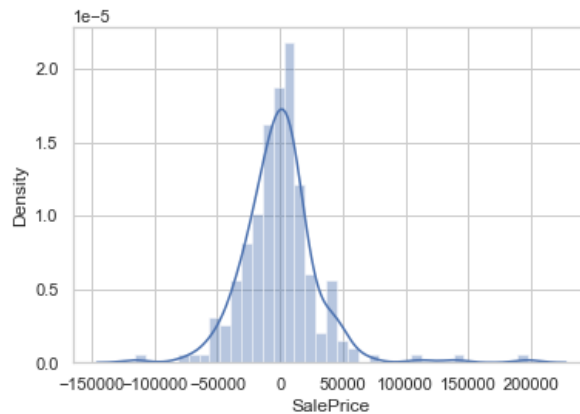
score=cross_val_score(RG,x,y,cv=10,scoring='r2')
print("Score:",score)
print('Mean Score:',score.mean())
print("Standard deviation:",score.std())

Score: [0.87093015 0.79113836 0.83668642 0.68392408 0.82132711 0.40663273
 0.78924226 0.76344277 0.66557158 0.87778786]
Mean Score: 0.7506683331134026
Standard deviation: 0.13280051158633877
```

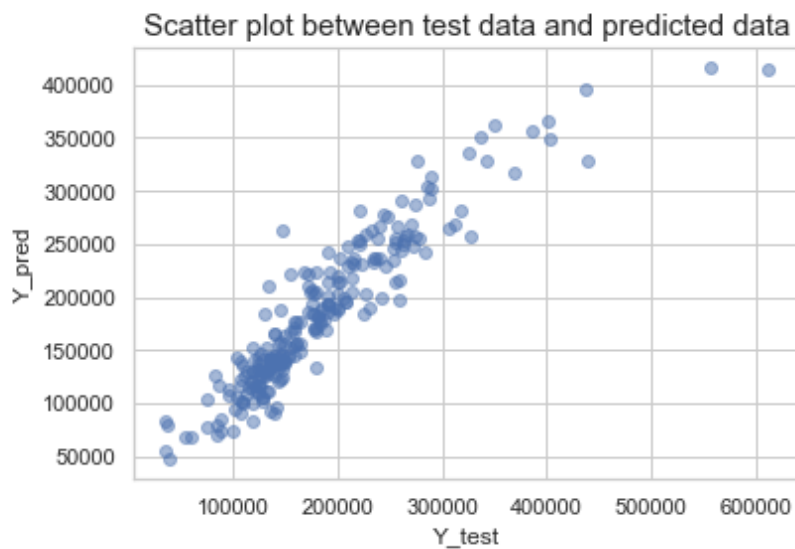
Let's plot the distribution plot and the Gaussian plot

```
sns.distplot(y_test-y_pred)
plt.show()
```

```
In [82]: # Let's plot the distribution plot and the Gaussian plot  
sns.distplot(y_test-y_pred)  
plt.show()
```



```
In [83]: # Let's plot the Scatter plot between test data and predicted data  
plt.scatter(y_test, y_pred, alpha=0.5)  
plt.xlabel("Y_test")  
plt.ylabel("Y_pred")  
plt.title("Scatter plot between test data and predicted data", fontsize=15)  
plt.show()
```



- Interpretation of the Results**

Data Pre-processing done by performing EDA (Exploratory Data Analysis), checking for best r2 score.

We will save our Model by Ridge Regression as it is giving us minimum Rmse score as it's having 30369.236527153855 and r2_score: 0.8594512207052254.

Model Saving

```
In [84]: #Ridge Regressor is giving us minimum Rmse score so we choose it as our final model.  
# Let's save our best model
```

```
import joblib  
joblib.dump(RG, 'Housing_Price_Project.pkl')
```

```
Out[84]: ['Housing_Price_Project.pkl']
```

```
In [85]: # Let's Load our save model
```

```
model=joblib.load('Housing_Price_Project.pkl')
```

```
In [86]: # Let's Test our save model
```

```
import sys  
nums= model.predict(x1)  
np.set_printoptions(threshold=sys.maxsize)  
print(nums)
```

```
5667824 88858855 483454 88866837 888886 88436846 458766 48877837
```

CONCLUSION

- **Key Findings and Conclusions of the Study**

Linear regression models assume that the relationship between a dependent continuous variable Y and one or more explanatory (independent) variables X is linear (that is, a straight line). It's used to predict values within a continuous range, (e.g. sales, price) rather than trying to classify them into categories (e.g. cat, dog).

- **Learning Outcomes of the Study in respect of Data Science**

- This dataset is Linear Regression in nature, we can verify data by using read method & get stats related information for each column using describe method.
- Visualizations, Pre-processing and Data Cleaning part was very crucial as without all these all method we were not able to judge the data effectively and won't be able to remove the outliers, handling null values and adding into the errors.
- Data contains numerical as well as categorical variable. So we handled them accordingly
- Check the r2 score using Mean absolute error, Mean squared error & get root mean squared error score.
- Train data using Linear Regression models to get the best score & finalise best score giver model for this dataset.
- Get the test score for same model.
- Save file using joblib library.

- **Limitations of this work and Scope for Future Work**

Visualizations helped a lot in finding out those outliers values and helped in finding out the features having direct relation between the feature and the label.

Its always good to to have complete data while performing model but 7-8 % of data can be excluded based on performance impact.

.